191870294-朱云佳-作业6

环境说明：**windows**使用**intellij**创建**maven**项目，环境配置与作业**5**相同，调试完毕后进入**WSL+hdfs**分布式环境运行

解题思路

标点、停词、大小写、数字和单词长度的处理沿用了作业5wordcount，此处不再赘述

Map

map输出的键值对是单词和所在文件名

```
if (!stoplist.contains(word.toString())&&(tmpword.length()>=3)){
    context.write(word, new Text(inputFileName));
    countmap+=1;
```

Reduce

reduce函数里定义hashmap类型的变量，用于存放输入reduce的特定单词在各文件中的出现次数。hashmap的键为文件名，值为出现次数。

统计次数

```
for (Text val : values) {
    if(!map.containsKey(val.toString())){
          System.out.println(val.toString());
        map.put(val.toString(),1);
    }
    else {
        map.put(val.toString(), map.get(val.toString()) + 1);
    }
}
```

若键不存在于map中，则创建并初始化值为1；否则更新已有的值

按词频排序

```
List<Map.Entry<String, Integer>> list = new ArrayList<>(map.entrySet());
Collections.sort(list, new Comparator<Map.Entry<String, Integer>>() {
            public int compare(Map.Entry<String, Integer> entry1, Map.Entry<String, Integer> entry2) {
                return entry2.getValue() - entry1.getValue();
            }
        }
);
StringBuilder docValueList = new StringBuilder();
for (int i=0;i<list.size();i++){
    docValueList.append(list.get(i).getKey()+"#"+list.get(i).getValue()+", ");
}
  context.write(key, new Text(String.valueOf(reducecount)+" "+docValueList.toString()));
context.write(key, new Text(docValueList.toString()));
```

map转entrySet之后，定义Collectionn.sort降序排列，最后用StringBuilder统一结果写入

实验结果

完整结果参见result文件夹part-r-00000，这里显示本地intellij和WSL分布式运行的结果

本地intellij

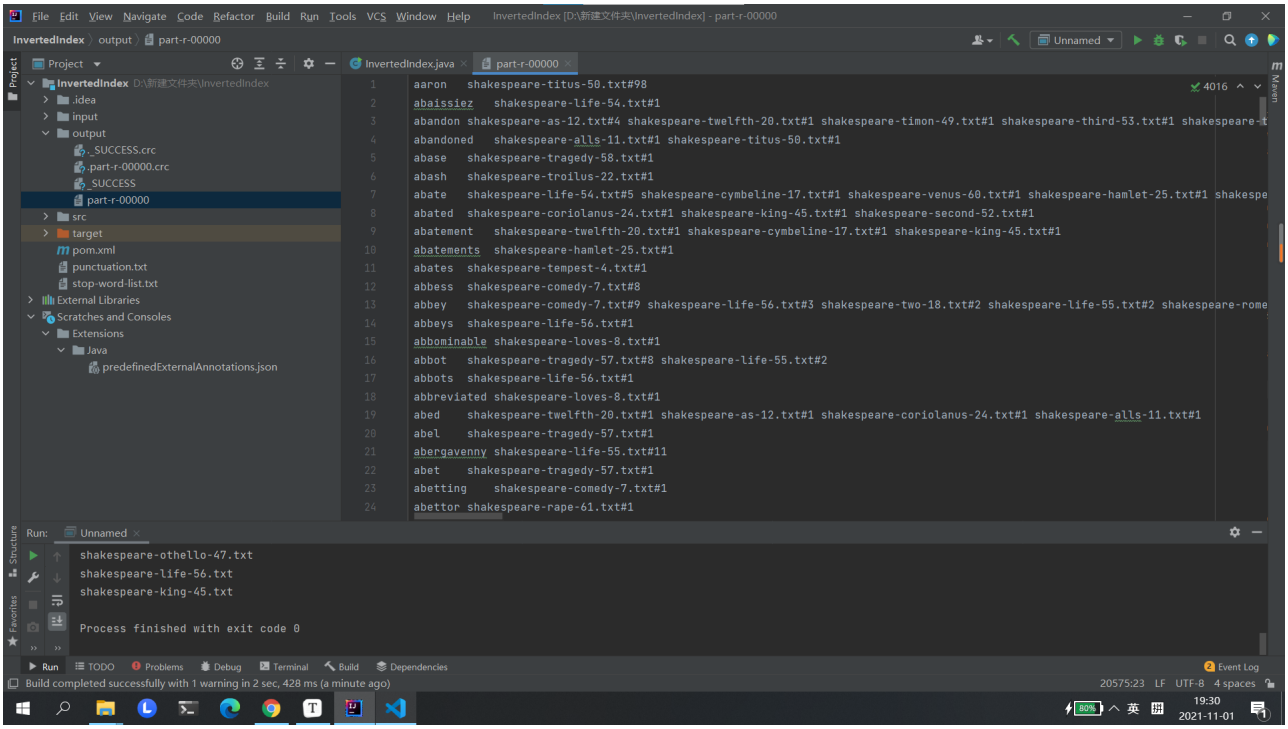

WSL分布式

用`hadoop jar`命令执行

```
zyj@LAPTOP-T1OKOQBM:~/hadoop/hadoop-3.3.0/bin$ hadoop jar /home/zyj/HW6/InvertedIndex-1.0-SNAPSHOT-jar-with-dependencies
.jar InvertedIndex /local/input output -skip /local/punctuation.txt
2021-11-01 20:04:59,955 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2021-11-01 20:05:00,298 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/
zyj/.staging/job_1635767942392_0001
2021-11-01 20:05:01,115 INFO input.FileInputFormat: Total input files to process : 40
2021-11-01 20:05:01,994 INFO mapreduce.JobSubmitter: number of splits:40
2021-11-01 20:05:02,083 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1635767942392_0001
2021-11-01 20:05:02,083 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-11-01 20:05:02,211 INFO conf.Configuration: resource-types.xml not found
2021-11-01 20:05:02,212 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2021-11-01 20:05:02,372 INFO impl.YarnClientImpl: Submitted application application_1635767942392_0001
2021-11-01 20:05:02,402 INFO mapreduce.Job: The url to track the job: http://LAPTOP-T1OKOQBM.localdomain:8088/proxy/appl
ication_1635767942392_0001/
2021-11-01 20:05:02,403 INFO mapreduce.Job: Running job: job_1635767942392_0001
2021-11-01 20:05:09,467 INFO mapreduce.Job: Job job_1635767942392_0001 running in uber mode : false
2021-11-01 20:05:09,467 INFO mapreduce.Job:  map 0% reduce 0%
2021-11-01 20:05:17,540 INFO mapreduce.Job:  map 22% reduce 0%
2021-11-01 20:05:24,619 INFO mapreduce.Job:  map 45% reduce 0%
2021-11-01 20:05:30,663 INFO mapreduce.Job:  map 52% reduce 0%
2021-11-01 20:05:31,669 INFO mapreduce.Job:  map 65% reduce 0%
2021-11-01 20:05:35,702 INFO mapreduce.Job:  map 75% reduce 0%
2021-11-01 20:05:36,706 INFO mapreduce.Job:  map 85% reduce 0%
2021-11-01 20:05:40,732 INFO mapreduce.Job:  map 100% reduce 0%
2021-11-01 20:05:41,737 INFO mapreduce.Job:  map 100% reduce 33%
2021-11-01 20:05:42,740 INFO mapreduce.Job:  map 100% reduce 100%
2021-11-01 20:05:44,756 INFO mapreduce.Job: Job job_1635767942392_0001 completed successfully
2021-11-01 20:05:44,813 INFO mapreduce.Job: Counters: 55
```

显示运行成功

```
zyj@LAPTOP-T1OKOQBM: ~/ha
        Total megabyte-milliseconds taken by all reduce tasks=35411968
    Map-Reduce Framework
        Map input records=158963
        Map output records=422310
        Map output bytes=13636813
        Map output materialized bytes=14481673
        Input split bytes=4947
        Combine input records=0
        Combine output records=0
        Reduce input groups=23596
        Reduce shuffle bytes=14481673
        Reduce input records=422310
        Reduce output records=23596
        Spilled Records=844620
        Shuffled Maps =40
        Failed Shuffles=0
        Merged Map outputs=40
        GC time elapsed (ms)=7694
        CPU time spent (ms)=87230
        Physical memory (bytes) snapshot=13907206144
        Virtual memory (bytes) snapshot=106673352704
        Total committed heap usage (bytes)=12883329024
        Peak Map Physical memory (bytes)=396607488
        Peak Map Virtual memory (bytes)=2610229248
        Peak Reduce Physical memory (bytes)=268935168
        Peak Reduce Virtual memory (bytes)=2615582720
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
```

发现hdfs系统上已生成结果文件

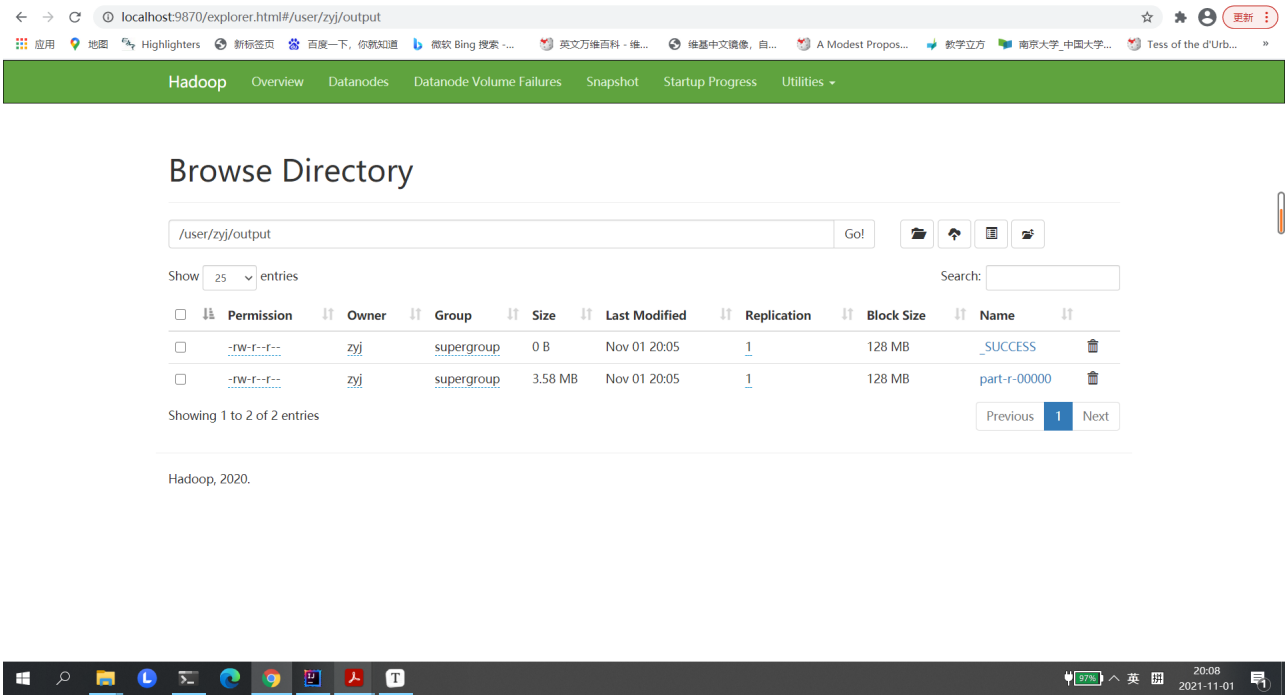

`cat`查看生成结果:

```
zyj@LAPTOP-T1OKOQBM:~/hadoop/hadoop-3.3.0/bin$ ./hdfs dfs -cat ./output/part-r-00000
```

(虽然由于行宽限制输出有点混乱但是)结果与本地一致!

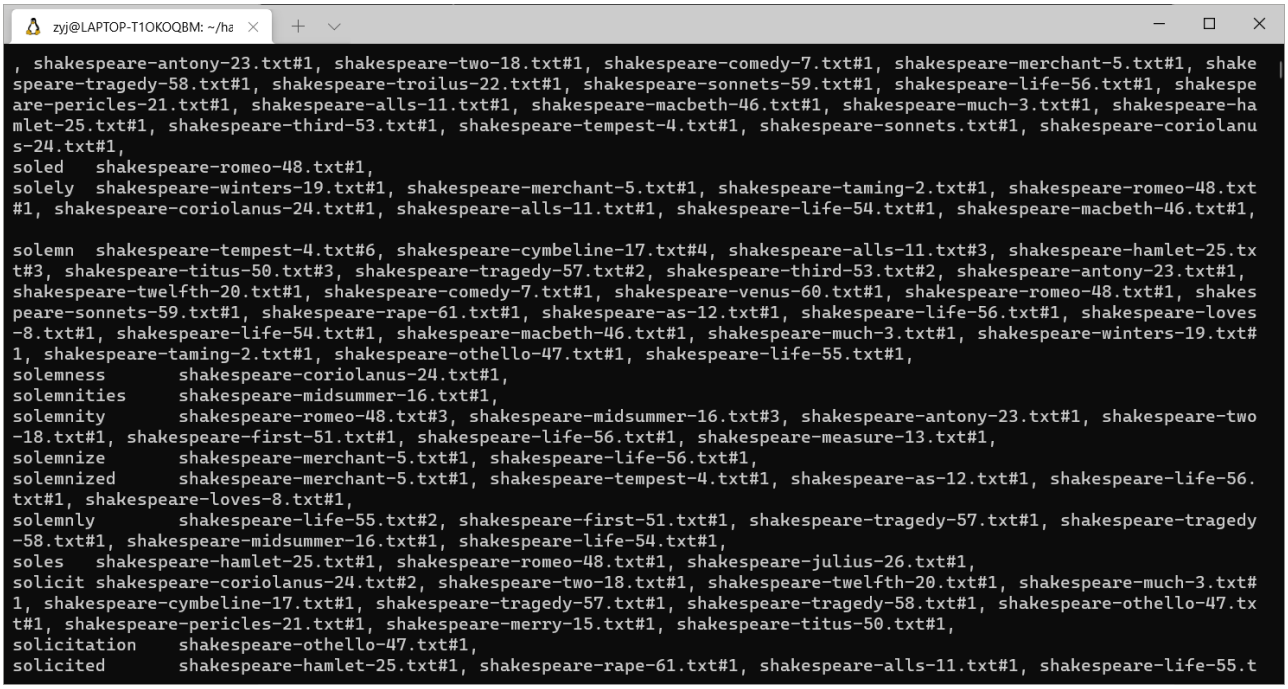

问题及解决方法

很快就完成了coding的框架，只是发现输出的时候：一是多余写入了"#1"，二是排序功能无法正常显示（后来推测与前者有关）

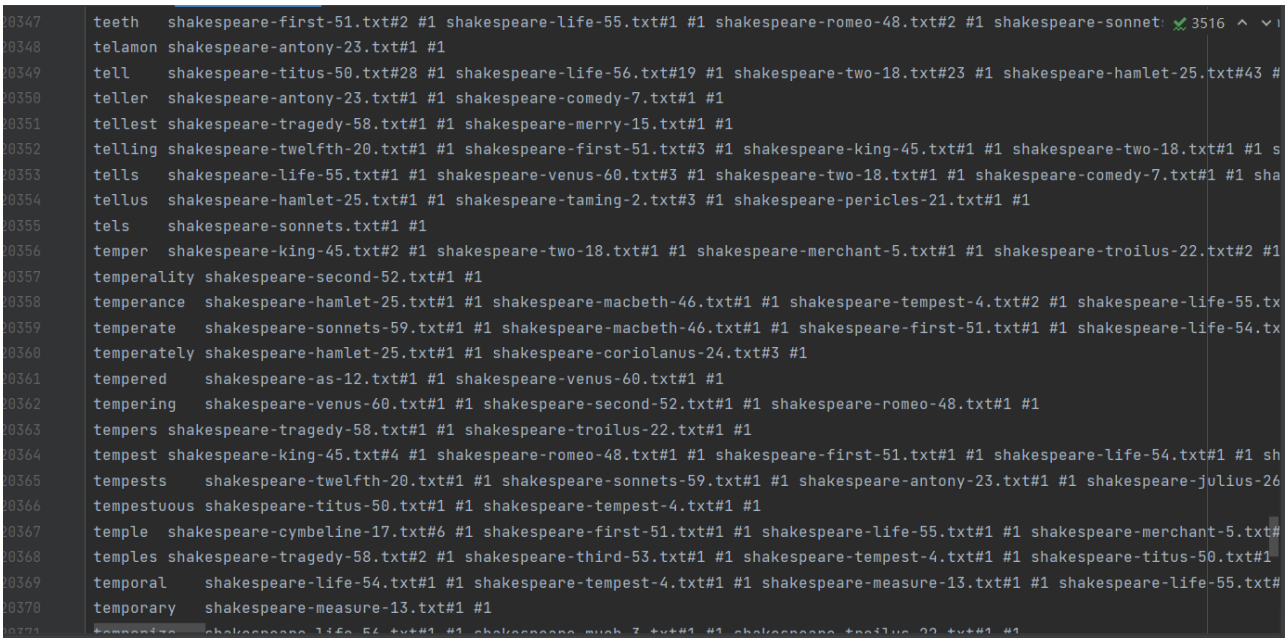

调整了context.write的输出（加入序列号）看起来，似乎是reduce一波之后，又来了一波：

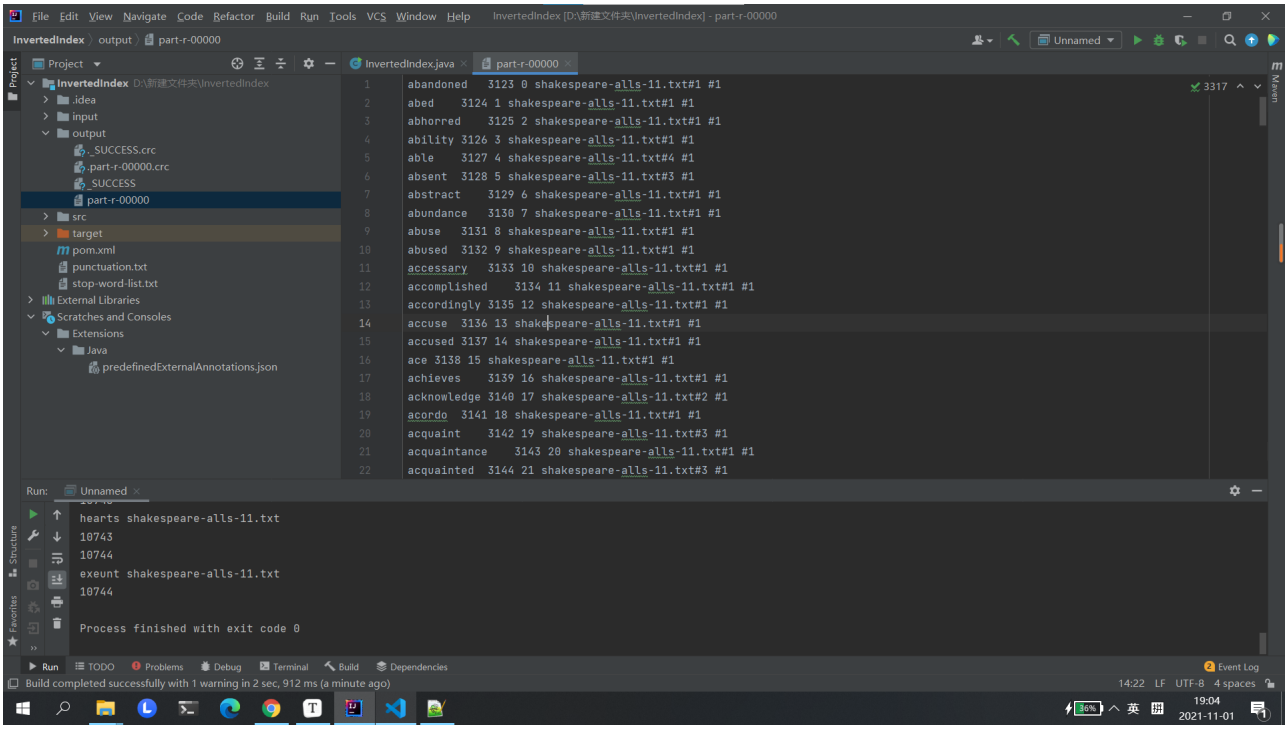

用println调试法--看起来它将之前的结果又统计了一遍（这里指的是shakespeare-alls-11.txt#2这种StringBuilder的内容，以下的print结果验证了猜想）

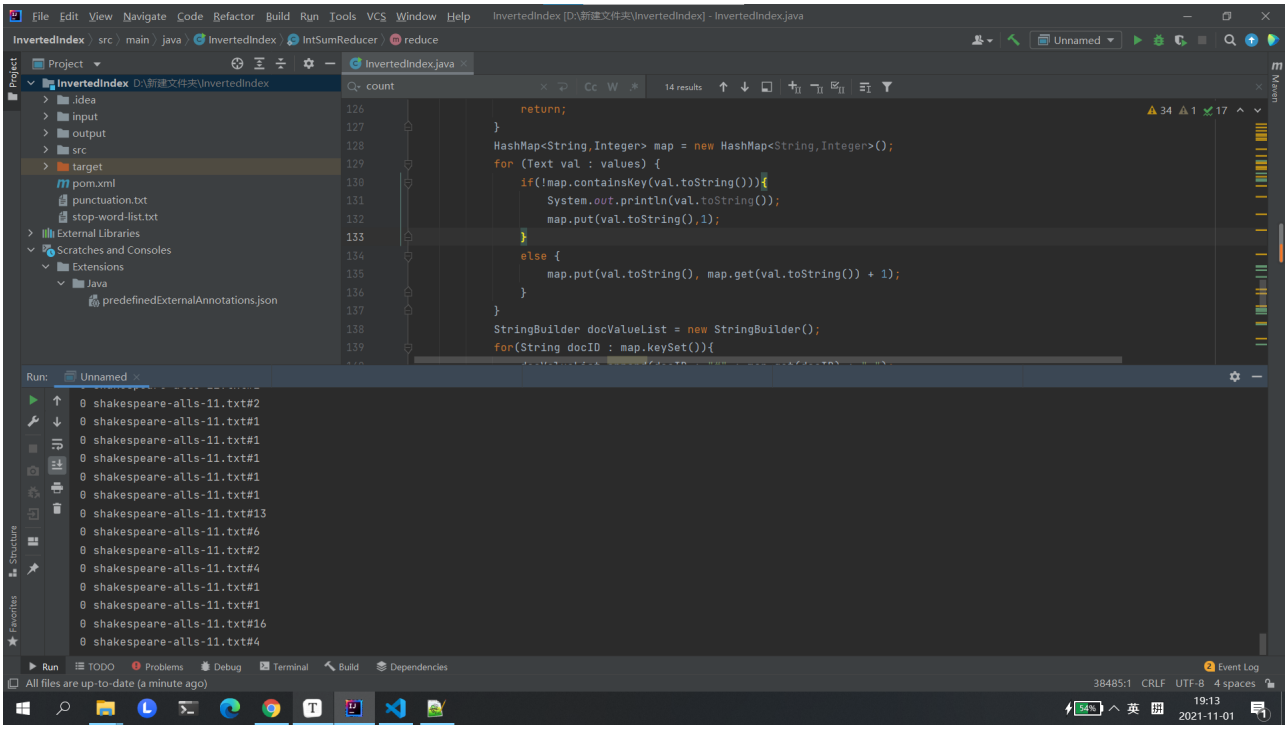

（其实花了9个小时断断续续调试才发现了上述端倪）...找来找去应该是combiner的锅（由于复用了wordcount的代码，没有去掉combiner，一开始也觉得不用去掉）

```
//        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(Text.class);
```

于是复习了combiner的作用（看来在这个情景下使用combiner会导致reduce的时候再统计一遍，因为第一波combiner过后value类型是Text）

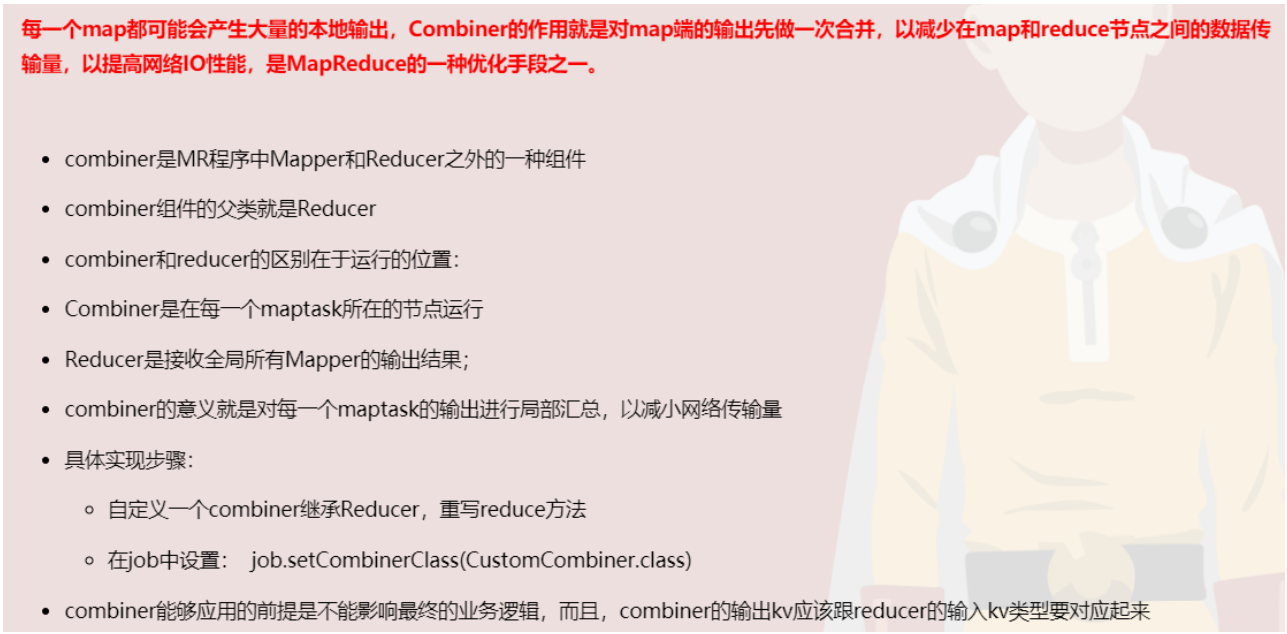

看来...要慎重考虑combiner的使用，解决方法就是注释掉combiner